

Introduction to Urban Data Science

CRP 4680/5680 (Spring 2025)



Week3 Data Management (II)

Wenzheng Li
Hazel (Yujin) Lee



OUTLINE

- 0. Quick Review
- 1. Data cleaning
 - Data Types
 - *NaN* Values
- 2. Joining multiple DataFrames
 - Concatenate Dataframes along the rows
 - Merging DataFrames along the columns
- 3. Slicing string columns





Review 0

df.columns

```
df_2012.columns
```

```
Index(['HouseID', 'CommunityID', 'TotalPrice', 'TransYear', 'Bedroom',  
      'Livingroom', 'Bathroom', 'Size', 'FloorLevel', 'WinSouth',  
      'WinSouthNorth', 'Decoration', 'TotalFloor', 'BuiltYear', 'Elevation',  
      'Heating', 'TransMonth', 'TransDay', 'District', 'CensusTract',  
      'XIAOQUWEB', 'SchQuality', 'NumSubway1km', 'Dist2Subway', 'HospQuality',  
      'Dist2Hosp', 'NumHosp1km', 'NumBus200m', 'Dist2CBD', 'Dist2Center',  
      'UnitPrice'],  
      dtype='object')
```

df_2012.columns returns an **index** object, not a list. To get the corresponding index based on value:

- `list(df_2012.columns).index("HouseID")`
- `df_2012.columns.tolist().index("HouseID")`



Indexing and Slicing a Dataframe

how to select a subset of a Dataframe?

- **Indexing**: simply selecting a particular row or column from a Dataframe.
- **Slicing**: selecting some rows and some columns
- Three ways of selecting particular rows and columns of a Dataframe
 - `df[]`
 - `df.loc[rows_label , columns_label]`
 - `df.iloc[row_position , column_position]`



- A **label**: one name in the column list or an index in the row index (the column at far left).
- A **position**: the corresponding position of column name or index in a sequence, starting from zero.

Label and Position

df											Column Positions
	0	1	2	3	4	5	6	7	8	9	Column names ... (Labels)
	HouseID	CommunityID	TotalPrice	TransYear	Bedroom	Livingroom	Bathroom	Size	FloorLevel	WinSouth	
0	0	BJCP85139027	1735	2080032.42	2012	3	2	2	124.62	3	1 ...
1	1	BJCY84814525	2023	1440023.27	2012	1	1	1	48.37	4	1 ...
2	2	BJHD61617745									1 ...
3	3	BJCY00382544									1 ...
4	4	BJCY84554915									1 ...

Row Positions Row index (Labels)

Task: select the first 2 rows and the first 2 columns

`df.loc[,]`

↑ ↑

row col

Labels Labels

`df.loc[0:1 , ['HouseID' , 'CommunityID']]`

`df.iloc[,]`

↑ ↑

row col

Positions Positions

`df.iloc[0:2 , 0:2]`

Filtering DataFrames

- `df.loc[df["Dist2Subway"] <= 1500, :] :`
 - step1, `df["Dist2Subway"] <= 1500` return a series with values of **False** or **True** (boolean type); this is known as **boolean indexing**
 - step2, it is enclosed by `df.loc[]` and can return a subset of the candidate rows
 - step3, assign the returned DataFrame to a new dataframe called `df_subway`
- In order to filter by more than one condition, you must:
 - Put all conditions in `()`
 - Separate the conditions by:
 - `|` if an **OR** condition
 - `&` if an **AND** condition

```
# houses within 1500m of subway stations with at least two Bathrooms  
df_2012.loc[(df_2012["Dist2Subway"] <= 1500) & (df_2012["Bathroom"] >= 2) ]
```

- Here, you cannot use the `.iloc` function because it used positional indexing. Boolean indexing (e.g., `(df_2012["Dist2Subway"] <= 1500) & (df_2012["Bathroom"] >= 2)`) is used to filter rows or columns based on conditions, and it works with `df.loc[]` or directly within square brackets `df[]`.

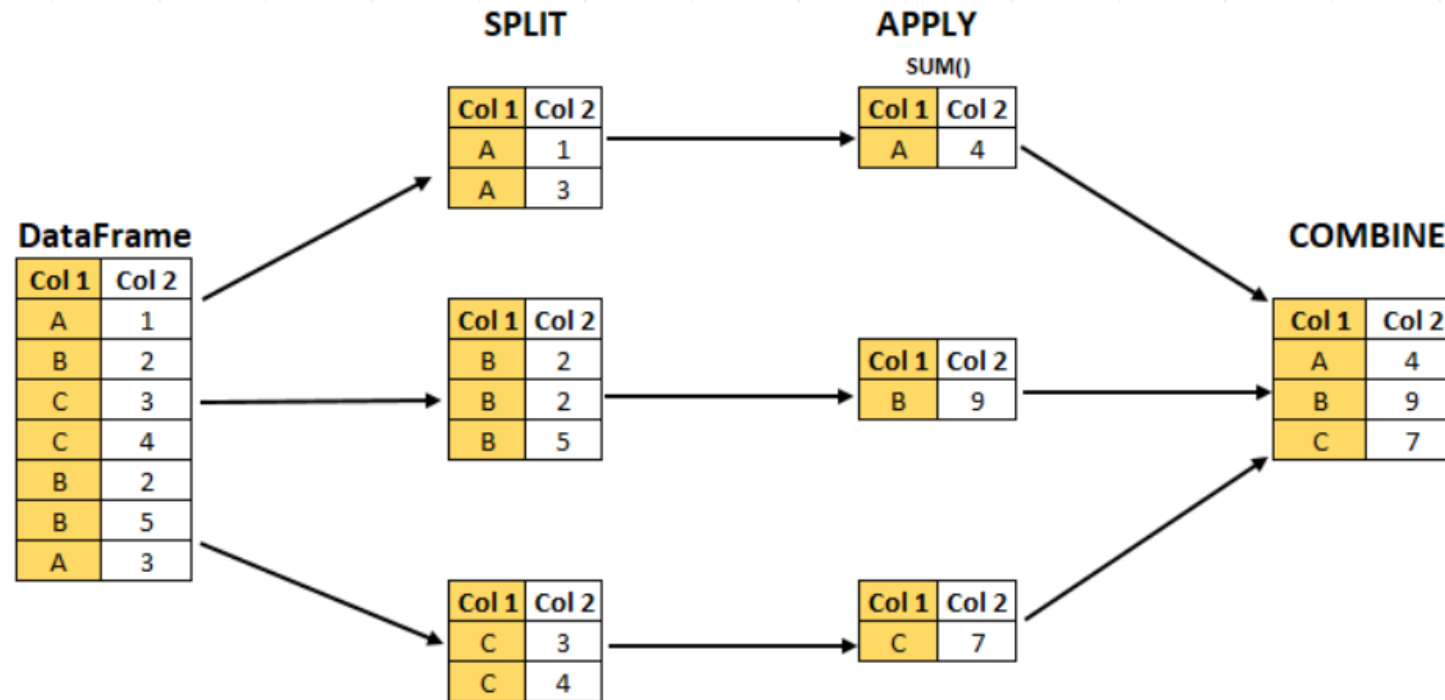
```
df_2012["Dist2Subway"] <= 1500
```

0	True
1	False
2	True
3	True
4	True
...	



By “group by” we are referring to a process involving one or more of the following steps:

- **Splitting** the data into groups based on some criteria.
- **Applying** a function to each group independently.
- **Combining** the results into a data structure.




```
# .reset_index()
df_2012.groupby("Sublevel")["UnitPrice"].mean()
```

UnitPrice

Sublevel	
Level 1	26882.939484
Level 2	25937.768385
Level 3	20309.530499

```
# .reset_index()
df_2012.groupby("Sublevel")["UnitPrice"].mean().reset_index()
```

	Sublevel	UnitPrice
0	Level 1	26882.939484
1	Level 2	25939.017190
2	Level 3	20328.317299

by default: **drop = False**


df.reset_index()

drop: If True, the current index is removed and is not added as a column

```
# .reset_index()
df_2012.groupby("Sublevel")["UnitPrice"].mean().reset_index(drop = True)
```

	UnitPrice
0	26882.939484
1	25937.768385
2	20309.530499

by default: **drop = True**



```
# groupby applied to UnitPrice
df_new = df_2012.groupby("Sublevel")["UnitPrice"].mean()
df_new
```

	UnitPrice
Sublevel	
Level 1	26882.939484
Level 2	25941.133662
Level 3	20328.317299

```
# .reset_index()
df_new.reset_index(inplace = False) # this is the default
df_new
```

	UnitPrice
Sublevel	
Level 1	26882.939484
Level 2	25941.133662
Level 3	20328.317299

↓

```
# .reset_index()
df_new2 = df_new.reset_index(inplace = False)
df_new2
```

	Sublevel	UnitPrice
0	Level 1	26882.939484
1	Level 2	25941.133662
2	Level 3	20328.317299

df.reset_index()

- **inplace=True**: Modifies the original DataFrame directly without returning a new DataFrame.
- **inplace=False**: Returns a new sorted DataFrame, leaving the original DataFrame unchanged.

```
# .reset_index()
df_new.reset_index(inplace = True)
df_new
```

	Sublevel	UnitPrice
0	Level 1	26882.939484
1	Level 2	25941.133662
2	Level 3	20328.317299




Data cleaning

1

Pandas dtype

- Each column/row in a Pandas (and GeoPandas) DataFrame has a data type, called *dtype* attribute.
- Here is the mapping between Pandas dtypes and python data types.
- Note that the object dtype means that the column is a mix of types or it's a string.


Pandas dtype	Python type	Usage
object	str or mixed	Text or mixed numeric and non-numeric values
int64	int	Integer numbers
float64	float	Floating point numbers
bool	bool	True/False values
datetime64	NA	Date and time values
timedelta[ns]	NA	Differences between two datetimes
category	NA	Finite list of text values



Null values

- Null values are when a particular value doesn't exist in a cell.
- In Pandas, you might see three different types of null values appear;
 - NaN (Not a Number), None, NA (only rarely)

	Column_None	Column_NaN	Column_String	Column_NA
0	1	1.1	apple	1
1	NaN	NaN	banana	NA
2	3	3.3	None	3
3	4	4.4	cherry	4



Null values

- **None** means a missing entry, but it's not a numeric type. It is of type **object** and is often found in columns that contain strings or mixed data types.
- **NaN** (Not a Number) used by Pandas for representing missing data in numeric columns.
- **Na** is Pandas' newer, more flexible missing data indicator that can be used across different data types.)

	Column_None	Column_NaN	Column_String	Column_NA
0	1	1.1	apple	1
1	NaN	NaN	banana	NA
2	3	3.3	None	3
3	4	4.4	cherry	4



NaN for Missing Value

- *NaN* is used for representing missing data in numeric columns.
- The data type of *NaN* **is float**, even if the rest of the column contains integers.
- To detect *NaN*, Pandas provides the *.isna()* and *.notna()* functions.
- Some Pandas operations will generate *NaN*. For example, when we concatenate or merge two DataFrames with a different number of columns or keys, the missing columns or rows will be filled with *NaN*.

[illegible]

NaN for Missing Value

- Removing data:
 - If it's an important cell, we might remove the entire row the cell belongs to.
- Imputing data:
 - We might want to replace it with:
 - The most frequent value (mode), if we think that there's some default value
 - The median value (if you think there are outliers in the sample that might be skewing the mean)
 - The average value (if you don't want the replaced data to influence your regression values).
 - Fill forward or backward: Fill missing values with the previous or next value (useful in time series data).
 - Or if you have more knowledge of the substantive topic (for ex: body temperature of mammals might typically be XX, but this species, it might be YY)
- Indicate that the data is missing in a new column
- Use linear regression or machine learning to predict the missing value.




Summary of Data Cleaning Steps and Exclusions

Table 5: Data Cleaning

			Total Observation: 356,942
Irrelevant Category			
Description	Count	Exclusion	
Buildings height < 3 stories	-696	1-story or 2-story buildings	
Non-flat unit	-1,534	dwellings with multiple levels, e.g., loft, duplex, split-level, and penthouse	
Not 70-year property	-8,002	40-year and 50-year properties, which are non-residential	
Not commodity property	-4,185	purchased public, affordable, state-owned, and all other non-commodity housing	
Not normal residential property	-2,901	villa, apartment, courtyard, bungalow, commercial-mixed	
Not basement	-954		
			Total Observation: 338,670
Outliers			
Description	Count	Exclusion	
Per square meter price	-446	< 2,000 RMB/square meter	
Total price	-1	< 100,000 RMB	
Floor area	-98	< 20 square meter or > 400 square meter	
			Total Observation: 338,125
Missing Value			
Description	Count	Exclusion	
Floor Level	-917	missing	
Elevator	-8,334	missing	
Year Built	-4,920	missing	
Heating	-3,181	missing	
			Total Observation: 320,773
Additional Drop			
Description	Count	Exclusion	
Outside 6th Ring Road	-16,232	observations outside the 6th ring road	
More than 5 bedrooms	-876	observations with more than 5 bedrooms	
More than 3 living rooms	-462	observations with more than 3 living rooms	
More than 4 bathrooms	-117	observations with more than 4 bathrooms	
Has no bathrooms	-112	observations with no bathrooms	
			Total Observation: 302,974

Li and Zhang (2024), Vertical Housing Gradients in Beijing. Under Review.





Joining multiple DataFrames **2**

Concatenating multiple Dataframes along the row axis (axis = 0)

- concatenating along the rows:
joining df2 to df1 vertically using `pd.concat(axis=0)`
- this means stacking your DataFrames on top of one another. If columns share the same names, they're combined into a single column; if not, new columns are created and filled with missing values.

df1					Result					
	A	B	C	D			A	B	C	D
0	A0	B0	C0	D0	x	0	A0	B0	C0	D0
1	A1	B1	C1	D1	x	1	A1	B1	C1	D1
2	A2	B2	C2	D2	x	2	A2	B2	C2	D2
3	A3	B3	C3	D3	x	3	A3	B3	C3	D3
df2					y	4	A4	B4	C4	D4
	A	B	C	D	y	5	A5	B5	C5	D5
4	A4	B4	C4	D4	y	6	A6	B6	C6	D6
5	A5	B5	C5	D5	y	7	A7	B7	C7	D7
6	A6	B6	C6	D6						
7	A7	B7	C7	D7						



2011

BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX CLASS AT PRESENT	BLOCK	LOT	EASE- MENT	BUILDING CLASS AT PRESENT	ADDRESS
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	372	19		C7	292 EAST THIRD STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	377	53		C2	269 EAST 7TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	385	53		C4	234 EAST 2ND STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	386	63		D7	215 EAST 3RD STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	400	53		C1	209 EAST 4 STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	402	28		C4	168 EAST 7TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	405	5		C1	182 AVENUE A
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	406	9		C7	506 EAST 13 STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	406	42		C7	543 EAST 12TH STREET
1	ALPHABET CITY	08 RENTALS - ELEVATOR APARTMENTS	2	379	53		D1	EAST 9TH STREET

df_2011

2012

BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX CLASS AT PRESENT	BLOCK	LOT	EASE- MENT	BUILDING CLASS AT PRESENT	ADDRESS
1	ALPHABET CITY	01 ONE FAMILY HOMES	1	372	38		S1	15 AVENUE D
1	ALPHABET CITY	01 ONE FAMILY HOMES	1	372	38		S1	15 AVENUE D
1	ALPHABET CITY	02 TWO FAMILY HOMES	1	376	32		S2	91 AVENUE D
1	ALPHABET CITY	03 THREE FAMILY HOMES	1	373	16		C0	326 EAST 4TH STREET
1	ALPHABET CITY	03 THREE FAMILY HOMES	1	393	9		C0	604 EAST 11TH STREET
1	ALPHABET CITY	04 TAX CLASS 1 CONDOS	1C	399	1102		R6	238 EAST 4TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	372	39		C7	11 AVENUE D
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	372	39		C7	11 AVENUE D
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	373	17		C3	328 EAST 4TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	385	38		C4	21-23 AVENUE C

df_2012



Combine Two DataFrames

BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX			BUILDING		
			CLASS AT	BLOCK	LOT	EASE- CLASS AT	PRESENT	ADDRESS
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	372	19	C7		292 EAST THIRD STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	377	53	C2		269 EAST 7TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	385	53	C4		234 EAST 2ND STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	386	63	D7		215 EAST 3RD STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	400	53	C1		209 EAST 4 STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	402	28	C4		168 EAST 7TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	405	5	C1		182 AVENUE A
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	406	9	C7		506 EAST 13 STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	406	42	C7		543 EAST 12TH STREET
1	ALPHABET CITY	08 RENTALS - ELEVATOR APARTMENTS	2	379	53	D1		EAST 9TH STREET
1	ALPHABET CITY	01 ONE FAMILY HOMES	1	372	38	S1		15 AVENUE D
1	ALPHABET CITY	01 ONE FAMILY HOMES	1	372	38	S1		15 AVENUE D
1	ALPHABET CITY	02 TWO FAMILY HOMES	1	376	32	S2		91 AVENUE D
1	ALPHABET CITY	03 THREE FAMILY HOMES	1	373	16	C0		326 EAST 4TH STREET
1	ALPHABET CITY	03 THREE FAMILY HOMES	1	393	9	C0		604 EAST 11TH STREET
1	ALPHABET CITY	04 TAX CLASS 1 CONDOS	1C	399	1102	R6		238 EAST 4TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	372	39	C7		11 AVENUE D
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	372	39	C7		11 AVENUE D
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	373	17	C3		328 EAST 4TH STREET
1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	385	38	C4		21-23 AVENUE C

2011

2012

df_combine

```
df_combine = pd.concat([df_2011, df_2012])
```

Combine Two DataFrames

	BOROUGH	NEIGHBORHOOD	BUILDINGCLASS	BLOCK	LOT	ADDRESS
0	1	ALPHABET CITY	07 RENTALS - WALKUP	372	19	292 EAST THIRD STREET
1	1	ALPHABET CITY	08 RENTALS - ELEVATOR	379	53	EAST 9TH STREET

df_2011

	BOROUGH	NEIGHBORHOOD	BUILDING_CLASS	BLOCK	LOT	PRESENT_ADDRESS
0	1	ALPHABET CITY	01 ONE FAMILY HOMES	372	38	15 AVENUE D
1	1	ALPHABET CITY	02 TWO FAMILY HOMES	376	32	91 AVENUE D

df_2012

	BOROUGH	NEIGHBORHOOD	BUILDINGCLASS	BLOCK	LOT	ADDRESS	BUILDING_CLASS	PRESENT_ADDRESS
0	1	ALPHABET CITY	07 RENTALS - WALKUP	372	19	292 EAST THIRD STREET	NaN	NaN
1	1	ALPHABET CITY	08 RENTALS - ELEVATOR	379	53	EAST 9TH STREET	NaN	NaN
0	1	ALPHABET CITY	NaN	372	38	NaN	01 ONE FAMILY HOMES	15 AVENUE D
1	1	ALPHABET CITY	NaN	376	32	NaN	02 TWO FAMILY HOMES	91 AVENUE D

df_combine

columns that don't match exactly become separate columns in the combined DF, leaving us with NaN values in places where the original DF had no corresponding data.



Combine Multiple DataFrames: Task

Let us combine the five housing datasets (2012 to 2016) together

initialization: df_combine ← Blank



Combine Multiple DataFrames: Task

Let us combine the five housing datasets (2012 to 2016) together

initialization: df_combine ← Blank

i = 2012: df_temp = housing data 2012

df_combine = df_temp

df_combine ← 2012



Combine Multiple DataFrames: Task

Let us combine the five housing datasets (2012 to 2016) together

initialization: df_combine ← Blank

i = 2012: df_temp = housing data 2012
df_combine = df_temp
df_combine ← 2012

i = 2013: df_temp = housing data 2013 2012
df_combine = pd.concat([df_combine, df_temp]) 2012
df_combine ← 2012, 2013



Combine Multiple DataFrames: Task


Let us combine the five housing datasets (2012 to 2016) together

initialization: df_combine ← Blank

i = 2012: df_temp = housing data 2012
df_combine = df_temp
df_combine ← 2012

i = 2013: df_temp = housing data 2013 2012
df_combine = pd.concat([df_combine, df_temp])
df_combine ← 2012, 2013

i = 2014: df_temp = housing data 2014 2012, 2013
df_combine = pd.concat([df_combine, df_temp])
df_combine ← 2012, 2013, 2014



Merging Dataframes along the column

- **Merging along the columns** means merging DF B to DF A horizontally based on a **merge key** (the column (or set of columns) whose values are used to match rows across the two DFs.).
- Function: `pd.merge()`
- `pd.concat()` can also be used to merge along columns by changing the argument `axis = 1`;
`pd.merge()` can **ONLY** be used to merge along the columns.


left				right				Result					
	key	A	B		key	C	D		key	A	B	C	D
0	K0	A0	B0	0	K0	C0	D0	0	K0	A0	B0	C0	D0
1	K1	A1	B1	1	K1	C1	D1	1	K1	A1	B1	C1	D1
2	K2	A2	B2	2	K2	C2	D2	2	K2	A2	B2	C2	D2
3	K3	A3	B3	3	K3	C3	D3	3	K3	A3	B3	C3	D3

Merging Dataframes along the column axis

- **Inner join** (`how='inner'`): Only rows with matching keys in both DataFrames are returned.
- **Outer join** (`how='outer'`): All rows from both DataFrames are returned, matched where possible, with missing values where not.
- **Left join** (`how='left'`): All rows from the left DataFrame are returned, matched with the right DataFrame where keys overlap, with missing values for unmatched right keys.
- **Right join** (`how='right'`): All rows from the right DataFrame are returned, matched with the left DataFrame where keys overlap, with missing values for unmatched left keys.

pandas.DataFrame.merge

```
DataFrame.merge(right, how='inner', on=None, left_on=None, right_on=None, left_index=False, right_index=False, sort=False, suffixes=('_x', '_y'), copy=None, indicator=False, validate=None)
```



See different merge scenarios at: https://pandas.pydata.org/docs/user_guide/merging.html

df_left		
	key	value_left
0	1	A
1	2	B
2	3	C



df_right		
	key	value_right
0	2	X
1	3	Y
2	4	Z

	key	value_left	value_right
0	2	B	X
1	3	C	Y

df_inner

	key	value_left	value_right
0	1	A	NaN
1	2	B	X
2	3	C	Y
3	4	NaN	Z

df_outer

	key	value_left	value_right
0	1	A	NaN
1	2	B	X
2	3	C	Y

df_left_join

	key	value_left	value_right
0	2	B	X
1	3	C	Y
2	4	NaN	Z

df_right_join



Q&A



df.shape

```
type(df_2012.shape)
```

tuple

```
T = ( 20, 'Jessa', 35.75, [30,60,90] )
```

T[0]

T[1]

T[2]

T[3]

tuple:

- a type of data structure very similar to **lists**
- a collection that is ordered and **unchangeable**.
- Tuple items are indexed. e.g., df_2012.shape[0], df_2012.shape[1]

